

AI for ABM: Move Up One Useful Level

A practical field guide to prompts, company context, reusable skills, governed automation, MCPs, and agent systems.

THE PROMISE

Pick one repeated ABM job. Make it grounded, inspectable, and reusable. Then move it up exactly one level.

Created for the ForgeX community by Trey Harnden

Uses the fictional North Peak AI Security and Alder Manufacturing demonstration.

Find your current level. Add the next layer.

0

Fancy Search

Broad questions and generalized answers

1

Contextual Copilot

One bounded task with supplied context

2

Grounded Workspace

Approved company and account context

3

Reusable Skill

A documented process and output contract

4

Governed Automation

A known workflow runs on a trigger

5

Goal-Driven Agent

The system chooses the plan and artifacts

Do not jump to Level 5. Earn the context, process, and controls underneath it.

Six terms that solve different problems.

PROMPT

The instruction and constraints for this run.

CONTEXT

The information the model can use for this job.

SKILL

A reusable operating manual with inputs, process, output, QA, and stop rules.

WORKFLOW

A known sequence of steps and handoffs.

MCP

A connection that gives an AI system approved access to resources and tools.

AGENT

A system that plans toward a goal, chooses tools, evaluates work, and stops at authority boundaries.

MCP gives access. It does not grant judgment, authority, or permission.

Build a company brain, not a document dump.

NORTH PEAK / CONTEXT

/brand-guidelines.md

/messaging-framework.md

/icp.md

/personas.md

/product-portfolio.md

/proof-and-claims.md

/source-manifest.md

SCOPED

Only what this job needs

SOURCE-BACKED

Claims retain provenance

FRESH

Owners and review dates

PERMISSION-AWARE

Access follows policy

VERSIONED

Changes stay inspectable

The model is replaceable. Context, skills, templates, tests, and decision history compound.

Port the method and blank templates. Leave employer-confidential content with the employer.

From an answer to an outcome.

1

Contextual Copilot

Use this seller and account context. Draft three campaign premises. Separate facts from hypotheses.

2

Grounded Workspace

Use the approved North Peak corpus and Alder evidence. Recommend one premise. Cite material claims.

3

Reusable Skill

Follow ABM Strategist. Produce the approval-ready brief through its process, output contract, and QA.

4

Governed Automation

Every Monday, refresh approved sources, update the brief, show the diff, and route exceptions. Do not publish.


5

Goal-Driven Agent

Increase qualified engagement. Decide which accounts need research, a brief, an experience, or no action. Ask before activation.

The big shift: Level 3 says produce an outcome through this process. Level 5 says here is the goal, decide what to produce.

A skill is a small operating manual for an agent.



ABM Strategist turns approved evidence into an inspectable, approval-ready account campaign brief.

USE THE SKILL TO IMPROVE THE SHARED PROCESS

Run multiple cases. Track corrections, missing inputs, exceptions, and unsupported claims. Fix the skill once so the next marketer starts stronger.

Human approval remains at targeting, claims, creative, activation, spend, and system-of-record writes.

Borrow the pattern. Keep the attribution.

MULTICHANNEL CAMPAIGN BUILDER | ERWANN LEFEVRE | LEVEL 3

Intake brief -> three-angle correction checkpoint -> channel rules -> complete output contract -> QA.

1:1 ABM ADS | IVAN FALCO | ADVANCED LEVEL 3 TO 4

SOPs and configs -> scripts -> DRAFT creation -> QA -> explicit activation gate. Draft-only. No activation or spend was tested.

PROPOSED PILOTS

Hiring verification | Signal interpretation | Audience ICP filtering | Dated company research

AUDIENCE ICP FILTER: CONDITIONAL PASS

22 synthetic records reconciled. Nine reached semantic review. Final: 12 match, 1 review, 6 no match, 3 excluded. Corrected regression suite: 45/45.

Reviewed list classification only. No enrichment, sending, CRM writes, or sequence enrollment.

External community work is MIT-licensed and creator-attributed. These are local evaluation references, not ForgeX certification or production adoption.

Ground the work in approved context.

WEEK 1

Choose one repeated job

List the company context people repeatedly paste, search for, or reconstruct.

WEEK 2

Build the corpus

Use approved brand, messaging, ICP, persona, product, proof, and source materials.

WEEK 3

Test three cases

Record contradictions, missing context, and generic or unsupported outputs.

WEEK 4

Assign ownership

Set owners and review dates, then place the corrected corpus in the approved workspace.

SUCCESS SIGNAL

Two marketers produce materially consistent work without rebuilding the company context.

Package the best-known process as a skill.

WEEK 1**Choose a bounded job**

Pick a repeated job with a clear start, finish, and review owner.

WEEK 2**Write the operating manual**

Define inputs, workflow, output contract, stop conditions, and quality checks.

WEEK 3**Run five cases**

Track corrections, exceptions, missing inputs, and unsupported claims.

WEEK 4**Improve the shared asset**

Add one template, one approved example, a named owner, and version history.

SUCCESS SIGNAL

Different marketers receive the same useful output structure with fewer corrections.

Run the proven workflow safely.

WEEK 1

Choose a safe trigger

Use one tested skill. Define the change and no-change cases.

WEEK 2

Map the controls

Document sources, permissions, state, update logic, retries, and approval points.

WEEK 3

Run draft-only

Compare every result with the manual process. Keep external action disabled.

WEEK 4

Activate with monitoring

Use an exception queue, audit record, and documented stop and recovery path.

SUCCESS SIGNAL

The workflow updates the right artifact, shows what changed, and routes exceptions without silent external action.

Measure outcomes, not AI activity.

SPEED

QUALITY

CONSISTENCY

CAMPAIGN VOLUME

PERSONALIZATION

PORTABLE-MARKETER CHECKLIST

- Context exports to readable files.
- Skills are not rewritten for one model.
- Material claims retain source and owner.
- Synthetic or approved examples test the process.
- A connector can change without losing business logic.
- Employer-confidential content stays with the employer.

LEVEL 4 TO 5

Use a one-to-one architecture session. Agent authority depends on your systems, data, permissions, goals, and failure tolerance.

Leave with one repeated job to move up.

- 1 Maturity rubric** Assess the current level and name the layer to add.
- 2 North Peak company brain** A fictional, reusable example of Level 2 grounding.
- 3 Context Corpus Builder** Turn approved brand, messaging, ICP, persona, product, and proof sources into a portable corpus.
- 4 ABM Strategist** Turn company and account evidence into an approval-ready brief.
- 5 30-day upgrade plans** Move from Level 1 to 2, 2 to 3, or 3 to 4.

REPOSITORY

github.com/0xTrey/forgex-practical-abm-skills

ForgeX will provide the attendee access method. Review every draft before external use.

Sources include Zapier's AI Fluency rubric, Anthropic's agent guidance, and creator-attributed Swan GTM Skills examples. Full URLs and evaluation boundaries are documented in the repository.